

Tutorial Agenda

Part 1: RISKS
1.5 hours

- Definitions 60 slides
 - Generative, Agentic, and Robotic AI
 - AI Safety and Security, AI Risk (threats, vulnerabilities, consequences/harms)
 - AI Confidentiality, Integrity, and Availability (CIA) and LLM, Agent, and Robot Failures
- Regulatory Environment
- LLM Safety and Security (biased, brittle, and baroque; harms, costs)
- Agent Safety and Security (belligerence)
- Robot Safety and Security

BREAK

Part 2: MITIGATIONS
1.5 hours

- AI Risk Modeling and Measurement 45 slides
- AI Governance Frameworks
- Risk Mitigation for LLMs, Agents and Robots (confidentiality, integrity, and availability), e.g., risk-based use cases, prompt guardrails, red teaming, human oversight.
- Future AI Research (ethics, resilience, privacy preservation, truth)

Take aways on [GitHub](#): Slides, References, Links to tools.